

S&DS 365 / 665
Intermediate Machine Learning

Course Wrap-up:

December 6

Yale

Reminders

- Assn 5 due tonight
- Final exam: Wed Dec 20 at 9am in HC L02
- Practice exams posted
- Review sessions (so far):
 - ▶ Tuesday, Dec 12, 7pm (Ruixiao)
 - ▶ Wednesday, Dec 13, 7pm (Zehao)

Last unit: Language/Sequence models

- Generative process, any sequence (of words, characters, stock prices, nucleotides...) is assigned a probability

$$p(x_1, \dots, x_n)$$

which can be factored as

$$p(x_1, \dots, x_n) = p(x_1)p(x_2 | x_1) \dots p(x_n | x_1, \dots, x_{n-1})$$

Transformers

The current state-of-the-art is based on *transformers*


- Attention is the key ingredient
- Rather than processing sequences word-by-word, transformers handle larger chunks of text at once
- Incorporates “interactions” between words and hidden states

Transformer architecture

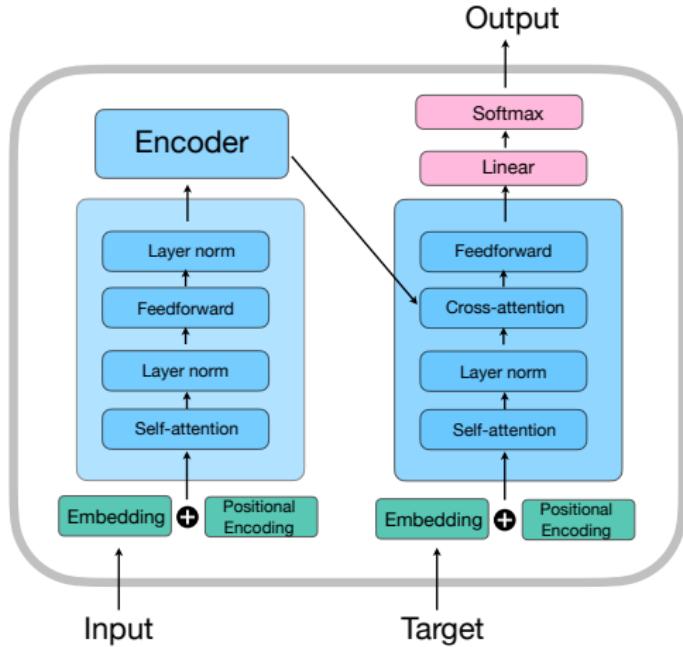
A Transformer is a seq2seq model based on encoder and decoder modules.

Transformers are powerful alternatives to RNNs that transform the encoder/decoder states using (multi-head) attention mechanisms.

Transformer architecture

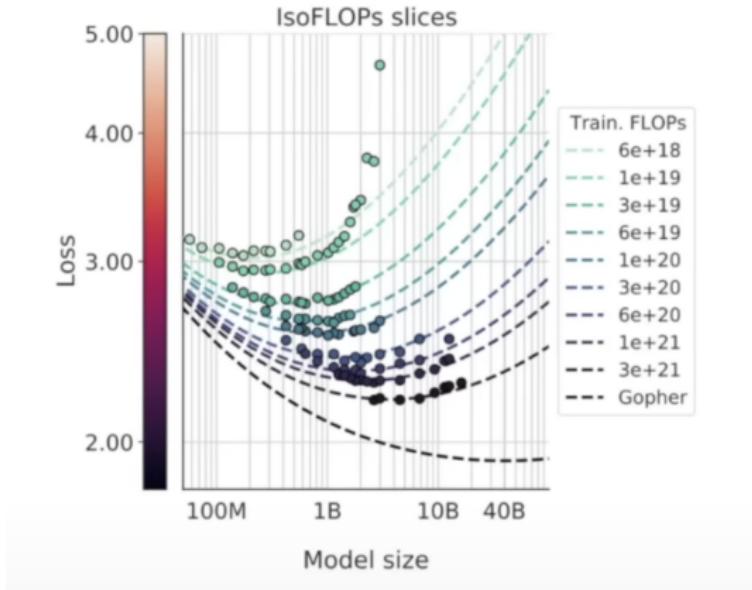


Transformer architecture



Two encoder layers and one decoder layer

LLM scaling laws: Bigger is better

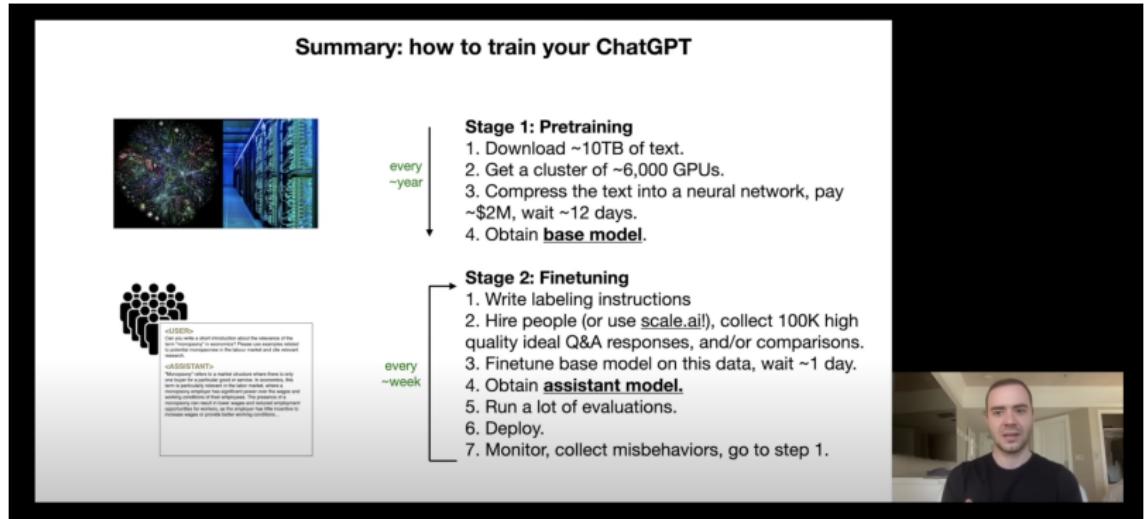


https://www.youtube.com/watch?v=zjkBMFhNj_g&t=25m40s

Finetuning (Training an LLM Decoder)

- An incredible amount of knowledge is stored implicitly in the weights of the transformer
- To be made useful, supervised learning and reinforcement learning are used to teach the model how to respond

Finetuning



https://www.youtube.com/watch?v=zjkBMFhNj_g&t=14m19s

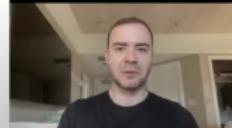


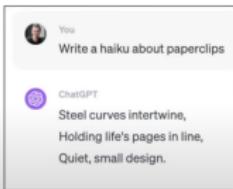
Finetuning

RLHF

The second kind of label: comparisons

It is often much easier to compare Answers instead of writing Answers.

Simple example: it's much easier to spot a good haiku than it is to generate one:



https://www.youtube.com/embed/zjkBMFhNj_g?start=1277&end=1340

Step 1

Collect demonstration data and train a supervised policy.

A prompt is sampled from our prompt dataset.

Explain reinforcement learning to a 6 year old.



We give treats and punishments to teach...



SFT



This data is used to fine-tune GPT-3.5 with supervised learning.

Step 2

Collect comparison data and train a reward model.

A prompt and several model outputs are sampled.

Explain reinforcement learning to a 6 year old.

A
It's reinforcement learning, the agent...
B
Explains rewards...
C
In machine learning...
D
We give treats and punishments to teach...



A labeler ranks the outputs from best to worst.

D > C > A > B

RM

D > C > A > B

This data is used to train our reward model.

Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.

A new prompt is sampled from the dataset.

Write a story about otters.



Once upon a time...

PPO



The PPO model is initialized from the supervised policy.


The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.

r_k

What's next?: Fast learning, slow thinking



AI struggles with IQ tests
Panther Media GmbH / Alamy Stock Photo

Always news—on both sides

The
Economist

≡ Menu

Weekly edition

The world in brief

Q Search ▾

Science and technology | Crystal balls

A Google AI has discovered 2.2m materials unknown to science

Zillions of possible crystals exist. AI can help catalogue them



Just one of trillions IMAGE: SCIENCE PHOTO LIBRARY

Nov 29th 2023

Save

Share

Give

Always news—on both sides

The
Economist

≡ Menu | Weekly edition | The world in brief | Q, Search ▾ |

Business | Of evils and evals

The world wants to regulate AI, but does not quite know how

There is disagreement over what is to be policed, how and by whom




IMAGE: MARI FOUZ

Oct 24th 2023 | BLETCHLEY PARK

Save

Share

Give

Advancing ML: Something for everyone

- Designing new methods
- Applying to new domains
- Solving deep mathematical puzzles
- Tackling unique engineering challenges
- Designing interfaces
- Collecting data
- Equitable use through policy and law
- Outreach and communication to broad communities

We've covered a lot of ground!

Calendar Fall 2023

Lectures: Monday/Wednesday 1:00-2:15pm

Luce Hall 101

Week	Dates	Topics	Demos & Tutorials	Lecture Slides	Readings & Notes	Assignments & Exams
1	Aug 30, Sep 1	Sparse regression	 Python elements  Pandas and regression  Lasso example	Wed: Course overview Fri: Sparse regression	PML Section 11.4 Notes on linear regression	
2	Sep 6	Smoothing and kernels	 Smoothing example  Using different kernels	Wed: Smoothing	PML Sections 16.3, 17.1 Notes on computing the lasso	Quiz 1
3	Sep 11, 13	Density estimation and Mercer kernels	 Density estimation demo  Mercer kernels (1/2)  Mercer kernels (2/2)	Mon: Density estimation Wed: Mercer kernels	Risk bounds for local smoothing Notes on Mercer kernels	 Assn 1 out

We've covered a lot of ground!

...						
4	Sep 18, 20	Neural networks and overparameterized models	CO np-complete example (1/2) CO np-complete example (2/2) TensorFlow playground	Mon: Neural networks Wed: Double descent	PML Sections 13.1, 13.2 Notes on backpropagation Notes on double descent	Quiz 2
5	Sep 25, 27	Convolutional neural networks	CO Convolution demo CO CNN demo	Mon: Double descent (continued) and Convolutional neural networks Wed: CNNs (continued)	PML Section 17.2 Notes on Bayesian inference Notes on nonparametric Bayes	Assn 1 in CO Assn 2 out
6	Oct 2, 4	Gaussian processes and approximate inference	CO Parametric Bayes CO Gaussian processes CO Gibbs sampling for image denoising	Mon: Gaussian processes Wed: Introduction to approximate inference	Notes on simulation	Quiz 3

We've covered a lot of ground!

7	Oct 9, 11	Variational inference	 Variational autoencoders	Mon: Variational inference Wed: VAEs	PML Section 20.3 Notes on variational inference	Assn 2 in  Assn 3 out
8	Oct 16	Midterm			Practice midterms	Oct 16: Midterm exam
9	Oct 23, 25	Graphs and structure learning	 Graphical lasso demo	Mon: Sparsity and graphs Wed: Discrete data and graph neural nets	Notes on graphs and structure learning Graph neural networks PML Section 23.4	
10	Oct 30, Nov 1	Deep reinforcement learning	 Q-learning demo  DQN demo	Mon: Reinforcement learning Wed: Deep reinforcement learning	Sutton and Barto, Section 6.5	Nov 1: Assn 3 in  Assn 4 out

We've covered a lot of ground!

Week						
11	Nov 6, 8	Policy gradient methods	CO Policy gradients demo CO Actor-critic demo	Mon: Policy gradient methods Wed: Actor-critic methods	Sutton and Barto, Section 13.1-13.3, 13.5	Quiz 4
12	Nov 13, 15	Sequential models	CO vanilla RNN CO Fakespeare GRU	Mon: HMMs and RNNs Wed: RNNs, GRUs, LSTMs, and all that	TensorFlow: Text generation Notes on HMMs and Kalman filters PML Chapter 15	Assn 4 in CO Assn 5 out
13	Nov 20, 22	No class, Thanksgiving break				
14	Nov 27, 29	Sequence-to-sequence models and Transformers	CO GPT-4 Python API	Mon: Sequence-to-sequence models Wed: Attention and transformers	Notes on mixtures PML Sections 15.4, 15.5	Quiz 5
15	Dec 4, 6	Transformers Societal issues	CO Transformer demo Minimal LLM decoder	Mon: Transformers, LLMs, and AI safety Wed: Course wrap up		Assn 5 in
17	Wed Dec 20, 9am, HQ L02	Final exam			Practice exams	Registrar: final exam schedule

Final exam

- Final exam Wednesday, Dec 20, 2023 at 9am in HQ L02
- <https://registrar.yale.edu/general-information/final-exams>
- Review sessions (see times/dates above)
- Length: About 1.5X Midterm
- Emphasis on material after midterm
- Cumulative, closed book, cheat-sheet

Your input

- Please complete a course review!
- I greatly value your comments and feedback
- Feel free to send me comments privately
- Let me know how you use and continue to learn ML!

Thank you!