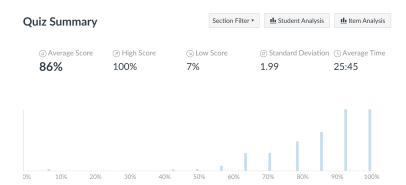
S&DS 365 / 665 Intermediate Machine Learning

Variational Inference

October 9

Reminders

- Assignment 2 due Wednesday
- Midterm week from today, in class
- Multiple review sessions scheduled (see Canvas)
- Quiz 3 scores posted; discussed in review sessions
- Practice midterms posted



Review sessions

For Today

- Variational inference: The ELBO
- Derivations (partial) and examples
- Variational autoencoders (VAEs)



Inverting generative models

Template for generative model:

- 1 Choose Z
- 2 Given z, generate (sample) X

We often want to invert this:

- Given x
- What is Z that generated it?

Inverting models

Bayesian setup:

- **1** Choose θ
- 2 Given θ , generate (sample) X

Posterior inference:

- Given x
- **2** What is θ that generated it?



Approximate inference

If we have a random vector $Z \sim p(Z \mid x)$, we might want to compute the following:

- marginal probabilities $\mathbb{P}(Z_i = z \mid x)$
- marginal means $\mathbb{E}(Z_i = z \mid x)$
- most probable assignments $z^* = \operatorname{arg\,max}_{Z} \mathbb{P}(\{Z_i = z_i\} \mid x)$
- maximum marginals $z_i^* = \arg \max_{z_i} \mathbb{P}(Z_i = z_i \mid x)$
- joint probability $\mathbb{P}(Z \mid x)$
- joint mean $\mathbb{E}(Z \mid x)$

Each of these quantities is intractable to calculate exactly, in general.

8

Variational methods

- Gibbs sampling is stochastic approximation
- Variational methods iteratively refine deterministic approximations
- Variational and Markov chain approximations originated in physics

Example 1: Interacting particles

We have a graph with edges E and vertices V. Each node i has a random variable Z_i that can be "up" (Z_i) or "down" $(Z_i = 0)$

$$\mathbb{P}_{\beta}(z_1,\ldots,z_n) \propto \exp\left(\sum_{s\in V} \beta_s z_s + \sum_{(s,t)\in E} \beta_{st} z_s z_t\right)$$

This is called an "Ising model" and is central to statistical physics.

Example 1: Interacting particles

We have a graph with edges E and vertices V. Each node i has a random variable Z_i that can be "up" (Z_i) or "down" $(Z_i = 0)$

$$\mathbb{P}_{\beta}(z_1,\ldots,z_n) \propto \exp\left(\sum_{s\in V} \beta_s z_s + \sum_{(s,t)\in E} \beta_{st} z_s z_t\right)$$

E are the set of edges, V are the vertices. Imagine the Z_i are votes of politicians, and the edges encode the social network of party affiliations

10

Stochastic approximation

Gibbs sampler

Iterate until converged:

- **1** Choose vertex $s \in V$ at random
- 2 Sample z_s holding others fixed

$$egin{aligned} heta_{ extbf{s}} &= \operatorname{sigmoid}\left(eta_{ extbf{s}} + \sum_{t \in extbf{N(s)}} eta_{st} z_{t}
ight) \ Z_{ extbf{s}} \mid heta_{ extbf{s}} \sim \operatorname{Bernoulli}(heta_{ extbf{s}}) \end{aligned}$$

You should verify that this is the correct conditional distribution of Z_s holding the others fixed.

 $sigmoid(u) = 1/(1 + e^{-u})$

1

Deterministic approximation

Mean field variational algorithm

Iterate until converged:

- **1** Choose vertex $s \in V$ at random
- 2 Update mean μ_s holding others fixed

$$\mu_{s} = \operatorname{sigmoid}\left(\beta_{s} + \sum_{t \in N(s)} \beta_{st} \, \mu_{t}\right)$$

Determinnistic vs. stochastic approximation

- The z_s variables are random
- The μ_s variables are deterministic
- The Gibbs sampler convergence is in distribution
- The mean field convergence is numerical
- The Gibbs sampler approximates the full distribution
- The mean field algorithm approximates the mean of each node

Example 2: A finite mixture model

Fix two distributions F_0 and F_1 , with densities $f_0(x)$ and $f_1(x)$, and form the mixture model

$$heta \sim \mathrm{Beta}(lpha, eta) \ X \mid heta \sim heta F_1 + (1 - heta) F_0.$$

The likelihood for data x_1, \ldots, x_n is

$$p(x_{1:n}) = \int_0^1 \text{Beta}(\theta \mid \alpha, \beta) \prod_{i=1}^n (\theta f_1(x_i) + (1-\theta) f_0(x_i)) d\theta.$$

Our goal is to approximate the posterior $p(\theta \mid x_{1:n})$

Stochastic approximation

Gibbs sampler

- **1** Sample $Z_i \mid \theta, x_{1:n}$ for i = 1, ..., n
- **2** Sample $\theta \mid z_{1:n}, x_{1:n}$

The first step is carried out by sampling

$$Z_i = \begin{cases} 1 & \text{with probability } \propto \theta f_1(x_i) \\ 0 & \text{with probability } \propto (1 - \theta) f_0(x_i) \end{cases}$$

15

Stochastic approximation

Gibbs sampler

- **1** Sample $Z_i \mid \theta, x_{1:n}$ for i = 1, ..., n
- **2** Sample $\theta \mid z_{1:n}, x_{1:n}$

The second step is carried out by sampling

$$\theta \sim \text{Beta}\left(\sum_{i=1}^{n} z_i + \alpha, n - \sum_{i=1}^{n} z_i + \beta\right).$$

Posterior over θ is approximated as *mixture* of Beta distributions; number of components is n+1

Variational inference: Strategy

- We'd like to compute $p(\theta, z | x)$, but it's too complicated.
- Strategy: Approximate as $q(\theta, z)$ that has a "nice" form
- q is a function of variational parameters, optimized for each x.
- Maximize a lower bound on p(x).

Variational inference: The ELBO

The ELBO is the following lower bound on $\log p(x)$:

$$\log p(x) = \int \sum_{z} q(z,\theta) \log p(x) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left(\frac{p(x,z,\theta) q(z,\theta)}{p(z,\theta \mid x) q(z,\theta)} \right) d\theta$$

$$= \sum_{z} \int q(z,\theta) \log \left(\frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta + \sum_{z} \int q(z,\theta) \log \left(\frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta$$

$$\geq \sum_{z} \int q(z,\theta) \log \left(\frac{p(x,z,\theta)}{q(z,\theta)} \right) d\theta$$

$$= H(q) + \mathbb{E}_{q}(\log p(x,z,\theta))$$

We maximize this over the parameters of q



Variational inference: The ELBO

The inequality above uses concavity of the logarithm:

$$\log\left(\sum_{\alpha} w_{\alpha} x_{\alpha}\right) \geq \sum_{\alpha} w_{\alpha} \log x_{\alpha}$$

So, if $q_{\alpha} \geq 0$ and $p_{\alpha} \geq 0$ sum (or integrate) to one, then

$$0 = \log \left(\sum_lpha oldsymbol{p}_lpha
ight) = \log \left(\sum_lpha oldsymbol{q}_lpha rac{oldsymbol{p}_lpha}{oldsymbol{q}_lpha}
ight) \geq \sum_lpha oldsymbol{q}_lpha \log \left(rac{oldsymbol{p}_lpha}{oldsymbol{q}_lpha}
ight)$$

Therefore

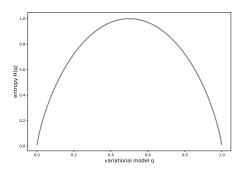
$$\sum_{lpha} q_lpha \log \left(rac{q_lpha}{p_lpha}
ight) \geq 0$$

18

Variational inference: The ELBO

The ELBO is $H(q) + \mathbb{E}_q(\log p)$

The entropy term H(q) encourages q to be spread out:



The cross-entropy $\mathbb{E}_q \log p$ tries to match q to p

Example 2: A finite mixture model

Fix two distributions F_0 and F_1 , with densities $f_0(x)$ and $f_1(x)$, and form the mixture model

$$heta \sim \mathrm{Beta}(lpha,eta) \ X \, | \, heta \sim heta F_1 + (1- heta) F_0.$$

The likelihood for data x_1, \ldots, x_n is

$$p(x_{1:n}) = \int_0^1 \text{Beta}(\theta \mid \alpha, \beta) \prod_{i=1}^n (\theta f_1(x_i) + (1-\theta) f_0(x_i)) d\theta.$$

Our goal is to approximate the posterior $p(\theta \mid x_{1:n})$

Variational approximation

Our variational approximation is

$$q(z,\theta) = q(\theta \mid \gamma_1, \gamma_2) \prod_{i=1}^n q_i^{z_i} (1-q_i)^{(1-z_i)}$$

where $q(\theta \mid \gamma_1, \gamma_2)$ is a Beta (γ_1, γ_2) distribution, and $0 \le q_i \le 1$ are n free parameters.

Need to maximize ELBO $H(q) + \mathbb{E}_q \log p$

Let's sketch part of the calculation

Variational approximation

First, we have

$$\log p(x,\theta,z) = \log p(\theta \mid \alpha,\beta) + \sum_{i=1}^{n} \left\{ \log \left(\theta^{z_i} f_1(x_i) \right) + \log \left(\theta^{1-z_i} f_0(x_i) \right) \right\}$$

Next we use identities such as

$$\mathbb{E}_q \log \theta = \psi(\gamma_1) - \psi(\gamma_1 + \gamma_2)$$

for the digamma function $\psi(\cdot)$.

After some calculus and algebra, we end up with the following algorithm (see the notes on the course web page for more detail)

Variational algorithm for mixture

Variational inference

Iterate the following steps for variational parameters $q_{1:n}$ and (γ_1, γ_2) :

• Holding q_i fixed, set $\gamma = (\gamma_1, \gamma_2)$ to

$$\gamma_1 = \alpha + \sum_{i=1}^n q_i$$
 $\gamma_2 = \beta + n - \sum_{i=1}^n q_i$

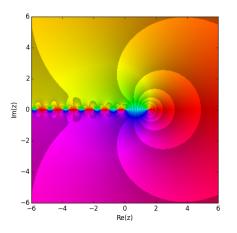
2 Holding γ_1 and γ_2 fixed, set q_i to

$$q_i = \frac{f_1(x_i) \exp \psi(\gamma_1)}{f_1(x_i) \exp \psi(\gamma_1) + f_0(x_i) \exp \psi(\gamma_2)}$$

After convergence, approximate posterior distribution over θ is

$$\widehat{p}(\theta \mid x_{1:n}) = \mathsf{Beta}(\theta \mid \gamma_1, \gamma_2)$$

Digamma function



 $\psi(x)$ is the digamma function https://en.wikipedia.org/wiki/Digamma_function

Deterministic approximation

- Convergence is numerical, not stochastic
- Posterior is approximated as a single Beta
- Very similar algorithm is used for topic models

Example 3: More general mixtures

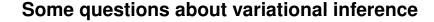
$$\theta \sim \mathsf{Dirichlet}(\alpha_1, \dots, \alpha_k)$$
 $X \mid \theta \sim \theta_1 F_1 + \dots + \theta_k F_k$

The likelihood for single data point x is

$$p(x) = \int \mathsf{Dirichlet}(\theta \,|\, \alpha_1, \dots, \alpha_k) \left(\sum_{j=1}^k \theta_j f_j(x) \right) \, d\theta.$$

When distributions F_j are learned, this is a "topic model." Variational inference is one of the most useful ways of training topic models

- Notes on variational methods: Please read Sections 1–4
- The other sections are more advanced / specialized material
- Next up: Variational autoencoders
- VAEs not on midterm, but please review basics of variational methods discussed so far, including the ELBO



Variational inference

Q: What is the best q we could use?

A: The true posterior $q(z, \theta \mid x) = p(z, \theta \mid x)$

Why? Because this maximizes the ELBO. Mathematically,

$$\sum_{z} \int q(z,\theta) \log \left(\frac{q(z,\theta)}{p(z,\theta \mid x)} \right) d\theta = 0$$

in this case, so the ELBO inequality is an equality.

Variational inference

Q: How does the ELBO regularize?

A: The entropy term favors distributions that are "spread out"

Why? For discrete distributions the maximum entropy distribution is uniform. For Gaussian, the entropy is $\log \sigma^2$ which favors σ^2 large.

Variational inference

Q: Is the ELBO easy to maximize?

A: No.

Why? In general it is non-convex, and the solution depends on where we start an iterative algorithm. This is unlike the Gibbs sampler, which converges to the right thing if we wait long enough.

Next class: Variational autoencoders

- Variational autoencoders are generative models that are trained using variational inference
- The "decoder" is a neural net that generates from a latent variable
- The "encoder" approximates the posterior distribution with another neural network trained using variational inference

Summary

- Gibbs sampling makes stochastic approximations
- Variational methods make deterministic approximations
- General recipe: Maximize ELBO over variational parameters
- Gives a powerful approach to generative modeling