


S&DS 365 / 665  
Intermediate Machine Learning

# Random Features and Double Descent

(continued)

September 25



Yale

# Reminders

- Assignment 1 is due Wednesday
- Assignment 2 posted Wednesday
- Quiz 2 grades posted

## Quiz Summary

Section Filter ▾

 Student Analysis

 Item Analysis

Ⓜ Average Score

**81%**

Ⓢ High Score

**100%**

Ⓣ Low Score


**0%**

Ⓢ Standard Deviation

**2.36**

⌚ Average Time

**18:24**




# Random features

Fix the weights at their random initializations, for all but the last layer

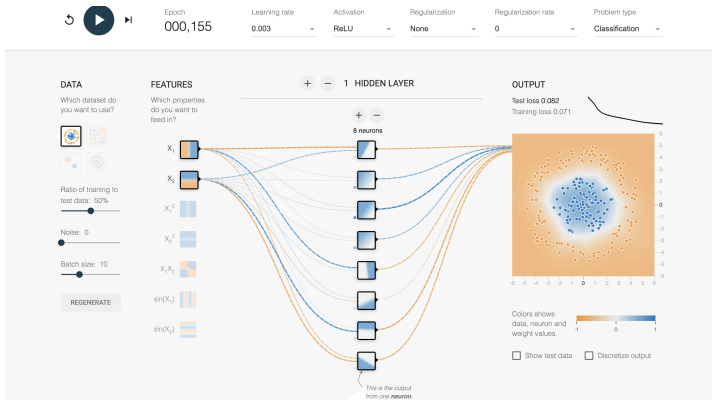
Just train the parameters  $\beta$

This is called the *random features model*. It's a linear model with random covariates obtained from the hidden neurons.

# Random features model



# Demo



<https://playground.tensorflow.org/>

# What's going on?

- These models are curiously robust to overfitting
- Why is this?
- Some insight: Kernels and double descent

# Double descent

We went over notes on minimum norm and double descent the board.

A problem on Assignment 2 will help solidify your understanding of this.



# OLS and minimal norm solution

OLS:  $p < n$

$$\hat{\beta} = (\mathbb{X}^T \mathbb{X})^{-1} \mathbb{X}^T \mathbf{Y}$$

Minimal norm solution:  $p > n$ :

$$\hat{\beta}_{\text{mn}} = \mathbb{X}^T (\mathbb{X} \mathbb{X}^T)^{-1} \mathbf{Y}$$

# “Ridgeless regression”

As  $\lambda$  decreases to zero, the ridge regression estimate:

- Converges to OLS in the “classical regime”  $\gamma < 1$
- Converges to  $\hat{\beta}_{mn}$  in “overparameterized regime”  $\gamma > 1$


# “Ridgeless regression”

This is a consequence of


$$(X^T X + \lambda I_p)^{-1} X^T = X^T (X X^T + \lambda I_n)^{-1}$$

which follows from the Woodbury formula (Assn 2).

# Double descent



# Double descent



# Double descent

- As  $\gamma \rightarrow \infty$ , the bias stays small
- Each entry of  $\frac{1}{p}\mathbf{X}\mathbf{X}^T$  is the average over an increasing number of identically distributed random vectors
- As a result, the variance decreases

# Neural tangent kernel

The *neural tangent kernel (NTK)* has been useful in understanding the performance of large neural networks, and the dynamics of stochastic gradient descent training.

# Parameterized functions

Suppose we have a parameterized function  $f_{\theta}(x) \equiv f(x; \theta)$

Almost all machine learning takes this form — for classification and regression, these give us estimates of the regression function

For neural nets, the parameters  $\theta$  are all of the weight matrices and bias (intercept) vectors across the layers.



# Feature maps

Suppose we have a parameterized function  $f_{\theta}(x) \equiv f(x; \theta)$

We then define a *feature map*

$$x \mapsto \varphi(x) = \nabla_{\theta} f(x; \theta_0) = \begin{pmatrix} \frac{\partial f(x; \theta_0)}{\partial \theta_1} \\ \frac{\partial f(x; \theta_0)}{\partial \theta_2} \\ \vdots \\ \frac{\partial f(x; \theta_0)}{\partial \theta_p} \end{pmatrix}$$

This defines a Mercer kernel

$$K(x, x') = \varphi(x)^T \varphi(x') = \nabla_{\theta} f(x; \theta_0)^T \nabla_{\theta} f(x'; \theta_0)$$

# Feature maps

This defines a Mercer kernel

$$K(x, x') = \varphi(x)^T \varphi(x') = \nabla_{\theta} f(x; \theta_0)^T \nabla_{\theta} f(x'; \theta_0)$$

*What is the NTK for the random features model?*

# Feature maps

The NTK for the random features model is

$$K(x, x') = h(x)^T h(x')$$

# Feature maps

*Conversely, a deep neural network with a large number of neurons is approximately equivalent to a random features model!*

Why?

# NTK and SGD

- The dynamics of stochastic gradient descent for deep networks has been studied
- Upshot: As the number of neurons in the layers grows, the parameters in the network barely change during training, even though the training error quickly decreases to zero

# NTK and random features

And, if the parameters only change by a small amount, a linear approximation can be used:

Let  $\theta = \theta_0 + \beta$ . Then

$$\begin{aligned} f(x, \theta) &\approx f(x, \theta_0) + \nabla_{\theta} f(x, \theta_0)^T \beta \\ &= \nabla_{\theta} f(x, \theta_0)^T \beta \end{aligned}$$

assuming that  $f(x, \theta_0) = 0$  (not a problem to assume this)

# NTK and random features

This tells us that the neural network is (approximately) equivalent to a random features model!

The random features are  $h(x) \equiv \nabla_{\theta} f(x, \theta_0)$

# Summary

- Neural nets are layered linear models with nonlinearities added
- Trained using stochastic gradient descent with backprop
- A surprise in the risk properties: Double descent
- Kernel connection: NTK